

CS2006: Advanced Programming Projects

Practical P2: Data analysis and visualisation Python

Project Introduction

This project required a group of three to analyze and visually present data on a provided census in Python using libraries like **pandas** and **matplotlib**.

The program should check for data inconsistencies and present the data in a way that can easily be interpreted.

Team Collaboration

For this project, we collaborated effectively as a team and managed to maintain our own programming style throughout the assignment.

Implementation

We successfully implemented all of the following tasks:

- **Basic**
- **Easy**
- **Medium**
- **Medium to Hard**
- **First Hard**: Displaying data using a map

Testing

We created an *invalidCensus.csv* file to test our cleanup process. When you change the file path in the cleanup.py file, you can see that it finds and removes any errors in the data, as well as pointing out some illogical data.

Problem Encountered

When we convert the notebook to an HTML file, the widgets are displayed, but they no longer function as there is no python backing it, so the plots are best viewed in notebook form.

One problem we encountered was trying to implement a script to convert our Python notebook to a PDF, which required downloading **LaTeX** with `sudo`, something we cannot do.

Instead, we used the following website to convert it:

html2pdf.com

Loading the data

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import ipywidgets as widgets
import seaborn as sns
from ipywidgets import interact, interactive_output, VBox, fixed
import numpy as np
import geopandas
import sys
sys.path.insert(0, '../code')
import labels
```

```
In [2]: %run "../code/cleanup.py"
df = pd.read_csv("../data/cleaned_data.csv")
```

=== Data Quality Check Results ===

Lenght of orignal dataset: 604351
Length of cleaned dataset: 604351

1. Missing Values:

resident_id_m	0
approx_social_grade	0
country_of_birth_3a	0
economic_activity_status_10m	0

```
ethnic_group_tb_6a      0
health_in_general       0
hh_families_type_6a     0
hours_per_week_worked   0
in_full_time_education  0
industry_10a            0
iol22cd                 0
legal_partnership_status_6a 0
occupation_10a          0
region                  0
religion_tb             0
residence_type          0
resident_age_7d         0
sex                     0
usual_short_student     0
dtype: int64
```

2. Duplicate Rows:
Number of duplicate rows: 0

3. Data Types:

```
resident_id_m           object
approx_social_grade     int64
country_of_birth_3a     int64
economic_activity_status_10m int64
ethnic_group_tb_6a      int64
health_in_general       int64
hh_families_type_6a     int64
hours_per_week_worked   int64
in_full_time_education  int64
industry_10a            int64
iol22cd                 object
legal_partnership_status_6a int64
occupation_10a          int64
region                  object
religion_tb             int64
residence_type          int64
resident_age_7d         int64
sex                     int64
usual_short_student     int64
dtype: object
```

4. Mutually Exclusive Values Rows:
None

5. Invalid Values Check:

Column: approx_social_grade
Number of invalid values: 0

Column: country_of_birth_3a
Number of invalid values: 0

Column: economic_activity_status_10m
Number of invalid values: 0

Column: ethnic_group_tb_6a
Number of invalid values: 0

Column: health_in_general
Number of invalid values: 0

Column: hh_families_type_6a
Number of invalid values: 0

Column: hours_per_week_worked
Number of invalid values: 0

Column: in_full_time_education
Number of invalid values: 0

Column: industry_10a
Number of invalid values: 0

Column: iol22cd
Number of invalid values: 0

Column: legal_partnership_status_6a
Number of invalid values: 0

Column: occupation_10a
Number of invalid values: 0

Column: region
Number of invalid values: 0

Column: religion_tb
Number of invalid values: 0

Column: residence_type
Number of invalid values: 0

Column: resident_age_7d
Number of invalid values: 0

Column: sex
Number of invalid values: 0

Column: usual_short_student
Number of invalid values: 0

Draw bar charts

Define dropdown box

```
In [3]: selection = widgets.Dropdown(
        options = [
            ("Social Grade", "approx_social_grade"),
            ("Country of Birth", "country_of_birth_3a"),
            ("Economic Activity", "economic_activity_status_10m"),
            ("Ethnic Group", "ethnic_group_tb_6a"),
            ("Health", "health_in_general"),
            ("Family", "hh_families_type_6a"),
            ("Hours Worked", "hours_per_week_worked"),
            ("Education", "in_full_time_education"),
            ("London", "iol22cd"),
            ("Partner", "legal_partnership_status_6a"),
            ("Region", "region"),
            ("Occupation", "occupation_10a"),
            ("Religion", "religion_tb"),
            ("Residence", "residence_type"),
            ("Age", "resident_age_7d"),
            ("Sex", "sex"),
            ("Usual Student", "usual_short_student"),
            ("Hours worked", "hours_per_week_worked"),
            ("Industry", "industry_10a"),
        ],
        value = "hours_per_week_worked",
        description = "Chart to display: ",
        style = {'description_width': 'initial'})
```

This dropdown box is used throughout the solution, mapping the easily readable variables to the names they're stored as in the csv file. Throughout the notebook, we also use a file that we created called *labels.py* that allows for transformations of numbers from the data set to what they actually represent.

Initialise bar chart functions

```
In [4]: def draw_bar_chart(column):
        df2 = df.copy()

        # Load labes from the labels module
        label = labels.label_map[column]
        title = labels.name_map[column] + " Distributions"
        xlabel = labels.name_map[column]

        # Count the number of records for the selected column
        counts = df2.groupby(column)[column].count()
        # Map the correct labels to the x axis variables
        counts.index = counts.index.map(label)
        counts.plot(kind='bar')

        # Set the title and labels
        plt.title(title)
        plt.xlabel(xlabel)
        plt.ylabel("Number of Residents")
```

By making the draw bar chart implimentation into a function it can take input from a dropdown widget, allowing it to display any record in the table.

Bar charts for basic requirements

```
In [5]: # Display the bar chart
out = interactive_output(draw_bar_chart, {'column': selection})
display(VBox([selection, out]))
```

To show the bar chart for age or occupation distribution as required in the basic requirements select either `Age` or `Occupation` from the dropdown box.

The age distribution shows us that for all sections other than under 15s and 65s and over there is around 70,000 records. It makes sense that the other two have more records since their age range is greater than the limited 10 year sections of the other 5.

Easy requirements

Initialise pie chart functions

```
In [6]: def draw_pie_chart(column):
        df2 = df.copy()

        # Load labels from the labels module
        label = labels.label_map[column]
        title = labels.name_map[column] + " Distributions"

        # Count the number of records for the selected column
        counts = df2.groupby(column)[column].count()
        # Map the correct labels to the variables
        counts.index = counts.index.map(label)
        counts.plot(kind='pie', autopct='%1.1f%%')

        # Set the title and labels
        plt.title(title)
        plt.ylabel("")
```

The same principal for the bar charts was applied to pie charts, allowing them to display any category again.

Pie Charts drawn

```
In [7]: # Display the pie chart
out = interactive_output(draw_pie_chart, {'column': selection})
display(VBox([selection, out]))
```

The easy requirements ask for pie charts for health descriptors and for ethnic groups, to see these select `Health` or `Ethnic Group` from the dropdown box.

By using the health section we can see that the majority of the records are in good health, with 48% in "Very good health" and 33% in "Good health", only 1.2% of the section have "Very bad health"

MEDIUM REQUIREMENTS

Group By tables for Hours Worked Vs Industry and Occupation Vs Social Grade

1. Number of records by hours by week and industry

```
In [8]: def draw_table(variable1, variable2):
        df2 = df.copy()

        # Only draw the table if the two variables are different
        if variable1 == variable2:
            print("Please select two different variables.")
            return

        # Load labels and display order from the labels module
        label1 = labels.label_map[variable1]
        label2 = labels.label_map[variable2]
        order1 = labels.order_map[variable1]
        order2 = labels.order_map[variable2]

        # Map the correct labels to the variables
        df2[variable1] = df2[variable1].map(label1)
        df2[variable2] = df2[variable2].map(label2)

        # Set the order of the categories
```

```

df2[variable1] = pd.Categorical(df2[variable1], categories=order1, ordered=True)
df2[variable2] = pd.Categorical(df2[variable2], categories=order2, ordered=True)

# Create a crosstab table
table = pd.crosstab(
    df2[variable1],
    df2[variable2],
    margins=True,
    margins_name="Total"
)

# Set the index and column names
table.index.name = labels.name_map[variable1]
table.columns.name = labels.name_map[variable2]

display(table)

```

Again the tables are designed to accept any two categories and display them accurately. The code for this section is more complicated however, due to the need to sort the table orders and label the variables for two different sections.

```

In [9]: # Define two new dropdowns for the table selection, using the same options
table_selection1 = widgets.Dropdown(
    options = selection.options,
    value = "hours_per_week_worked",
    description = "Table to display: ",
    style = {'description_width': 'initial'}
)
table_selection2 = widgets.Dropdown(
    options = selection.options,
    value = "industry_10a",
    description = "Table to display: ",
    style = {'description_width': 'initial'}
)

# Display the table
out = interactive_output(draw_table, {
    'variable1': table_selection1,
    'variable2': table_selection2
})

display(VBox([table_selection1, table_selection2, out]))

```

The two required tables for this section are mapping hours worked per week against industry and occupation against approximate social grade, again selectable by using the dropdown widget.

Analysis of the first and we can see that there is 171,000 records in "does not apply for" both records, presumably people who don't work. More interestingly are the 150,000 people who work in an industry both chose "Does not apply", these people could potentially be working very variable hours per week or currently be on leave.

When looking at the occupation against approximate social grade table we can see the expected results that the higher social grade records typically work in professional occupations or as managers and directors while the lower social grades are far more likely to work in Elementary Occupations

2. Number of records by occupation and approximate social grade

Extra addition to the above - Heatmaps for both tables

Heatmap for Hours Worked vs Industry

```

In [10]: def draw_heatmap(variable1, variable2, colour):
    df2 = df.copy()

    # Only draw the heatmap if the two variables are different
    if variable1 == variable2:
        print("Please select two different variables.")
        return

    # Load labels and display order from the labels module
    label1 = labels.label_map[variable1]
    label2 = labels.label_map[variable2]
    order1 = labels.order_map[variable1]
    order2 = labels.order_map[variable2]

    # Map the correct labels to the variables
    df2[variable1] = df2[variable1].map(label1)
    df2[variable2] = df2[variable2].map(label2)

    # Set the order of the categories

```

```

df2[variable1] = pd.Categorical(df2[variable1], categories=order1, ordered=True)
df2[variable2] = pd.Categorical(df2[variable2], categories=order2, ordered=True)

# Create a heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(
    pd.crosstab(df2[variable1], df2[variable2]),
    annot=True, fmt='d', cmap=colour, cbar=False
)

# Set the title and labels
plt.title(labels.name_map[variable1] + " vs " + labels.name_map[variable2])
plt.ylabel(labels.name_map[variable1])
plt.xlabel(labels.name_map[variable2])

```

The Heatmap function looks nearly identical to the previous table function, with very similar labelling and drawing. The key change is the use of the sns heatmap function rather than the pandas crosstab

```

In [11]: # Define a widget for the colour selection
colour_selection = widgets.Dropdown(
    options = ["Blues", "Greens", "Reds", "Purples", "Oranges"],
    description = "Colour scheme: ",
    style = {'description_width': 'initial'}
)

# Display the heatmap
out = interactive_output(draw_heatmap, {
    'variable1': table_selection1,
    'variable2': table_selection2,
    'colour': colour_selection
})

display(VBox([table_selection1, table_selection2, colour_selection, out]))

```

By drawing the same required tables from above in a heatmap the patterns previously noted are even more easy to see.

Medium Requirments 2

Economic Activitiy Analysis

```

In [12]: def draw_economic_bar_chart(column):
    df2 = df.copy()

    # Ensure the selected column is not "economic_activity_status_10m"
    if column == "economic_activity_status_10m":
        print("Please select a different variable.")
        return

    # Load labels from the labels module
    label = labels.label_map[column]
    xlabel = labels.name_map[column]

    # Create a new dataframe with only the relevant records
    economic_active = df2[df2['economic_activity_status_10m'].isin([1, 2, 3, 4])] # Active status codes 1,2,3,4
    economic_inactive = df2[df2['economic_activity_status_10m'].isin([5, 6, 7, 8, 9])] # Inactive status codes

    # Count the number of records for the selected column
    active_by_variable = economic_active.groupby(column)[column].count()
    inactive_by_variable = economic_inactive.groupby(column)[column].count()

    # Map the correct labels to the x axis variables
    active_by_variable.index = active_by_variable.index.map(label)
    inactive_by_variable.index = inactive_by_variable.index.map(label)

    x = np.arange(len(active_by_variable.index))
    w = 0.4

    # Create a bar chart
    fig, ax = plt.subplots(figsize=(10, 6))
    ax.bar(2*x + w/2, active_by_variable.values, w, label='Active', alpha=0.7)
    ax.bar(2*x - w/2, inactive_by_variable.values, w, label='Inactive', alpha=0.7)
    plt.title('Economically Active vs Inactive People by ' + xlabel)
    plt.xticks(2*x, active_by_variable.index, rotation=90)
    plt.legend()
    plt.xlabel(xlabel)
    plt.ylabel('Count')

```

```
plt.show()
```

The same repeatable principle was used on the comparison bar chart function. This time with the caveat that two bar plots are being drawn onto the same graph and that a new dataframe has to be created for both economically active and inactive, using the economic activity status.

```
In [13]: # Display the bar chart
out = interactive_output(draw_economic_bar_chart, {
    'column': selection
})

display(VBox([selection, out]))
```

Using this chart you can compare the number of economically active and inactive people based on both age and health groups. When analysing the age charts as expected the least active group is the 65 and over group, due to the large sections of this group which will be retired. The youngest age group also has a significant inactive section, this time due to the student population. Interestingly it's clear that the sum of both active and inactive people for the oldest group is more than that of any other group, this is due to both the larger age range of this group as well as the removal of "Does not apply" having the smallest effect on this section.

Comparing health with economic activity again produces a very expected value, with the majority of people in good and very good health economically active, while people in bad or very bad health are far less likely to be active.

```
In [14]: def draw_student_economic_bar_chart(column):
    df2 = df.copy()

    if column == "economic_activity_status_10m":
        print("Please select a different variable.")
        return

    # Load labels from the labels module
    label = labels.label_map[column]
    xlabel = labels.name_map[column]

    # Create a new dataframe with only the relevant records
    student = df2[df2['economic_activity_status_10m'].isin([4, 6])]
    not_student = df2[df2['economic_activity_status_10m'].isin([1, 2, 3, 5, 7, 8, 9])]

    # Count the number of records for the selected column
    student_for_variable = student.groupby(column)[column].count()
    not_student_for_variable = not_student.groupby(column)[column].count()

    # Map the correct labels to the x axis variables
    student_for_variable.index = student_for_variable.index.map(label)
    not_student_for_variable.index = not_student_for_variable.index.map(label)

    x = np.arange(len(student_for_variable.index))
    w = 0.4

    # Create a bar chart
    fig, ax = plt.subplots(figsize=(10, 6))
    ax.bar(2*x + w/2, student_for_variable.values, w, label='Students')
    ax.bar(2*x - w/2, not_student_for_variable.values, w, label='Non-Students')
    plt.title(xlabel + ' for Students vs Non-Students')
    plt.xticks(2*x, student_for_variable.index, rotation=90)
    plt.legend()
    plt.xlabel(xlabel)
    plt.ylabel('Count')
    plt.show()
```

```
In [15]: # Display the bar chart
out = interactive_output(draw_student_economic_bar_chart, {
    'column': selection
})

display(VBox([selection, out]))
```

Hard Requirement

Map of the Uk

Our first goal is to create a map of the UK that shows the amount of records that come from each region First we must import the shapefile that will give us the geographic data. We are using the geopandas library

```
In [16]: df = pd.read_csv("../data/publicmicrodatateachingsample.csv")
```

```
map = geopandas.read_file("../data/shapeFileData/NUTS1_Jan_2018_UGCB_in_the_UK.shp")
map.head()
```

Out[16]:

	nuts118cd	nuts118nm	bng_e	bng_n	long	lat	GlobalID	geometry
0	UKC	North East (England)	417313	600358	-1.72890	55.297031	628608f5-673c-45c8-9082-f57ed9604a20	MULTIPOLYGON (((397942.931 657535.008, 400000 ...
1	UKD	North West (England)	350015	506280	-2.77237	54.449451	0413c215-94f7-4f2e-a0fd-37d517177413	MULTIPOLYGON (((357136.203 587050.898, 356937....
2	UKE	Yorkshire and The Humber	446903	448736	-1.28712	53.932640	d608863c-86bd-4877-b3b8-247794a0825a	POLYGON ((479999.978 517670.16, 481340.496 515...
3	UKF	East Midlands (England)	477660	322635	-0.84967	52.795719	428c9691-f97b-4fa3-aef8-b3d8c0789cf9	POLYGON ((516022.702 412210.895, 515854.099 41...
4	UKG	West Midlands (England)	386294	295477	-2.20358	52.556969	a8a5953c-6a23-4fca-b1c7-96d80a29dd95	POLYGON ((409402.997 365710.796, 412633.099 36...

However, we can see that rather than using the region codes that the census data uses, this data set uses *nuts118cd* as the region identifier. To fix this, we will map these to the region codes that our data set uses. We will also rename the column name to *region* to make it more user readable.

In [17]:

```
# Rename the region codes to the modern version that the census data uses
map["nuts118cd"] = map["nuts118cd"].map({"UKC": "E12000001", "UKD": "E12000002", "UKE": "E12000003", "UKF": "E12000004", "UKG": "E12000005", "UKH": "E12000006", "UKJ": "E12000007", "UKK": "E12000008", "UKL": "E12000009", "UKM": "E12000010", "UKN": "E12000011", "UKO": "E12000012", "UKP": "E12000013", "UKQ": "E12000014", "UKR": "E12000015", "UKS": "E12000016", "UKT": "E12000017", "UKU": "E12000018", "UKV": "E12000019", "UKW": "E12000020", "UKX": "E12000021", "UKY": "E12000022", "UKZ": "E12000023", "UKAA": "E12000024", "UKAB": "E12000025", "UKAC": "E12000026", "UKAD": "E12000027", "UKAE": "E12000028", "UKAF": "E12000029", "UKAG": "E12000030", "UKAH": "E12000031", "UKAI": "E12000032", "UKAJ": "E12000033", "UKAK": "E12000034", "UKAL": "E12000035", "UKAM": "E12000036", "UKAN": "E12000037", "UKAO": "E12000038", "UKAP": "E12000039", "UKAQ": "E12000040", "UKAR": "E12000041", "UKAS": "E12000042", "UKAT": "E12000043", "UKAU": "E12000044", "UKAV": "E12000045", "UKAW": "E12000046", "UKAX": "E12000047", "UKAY": "E12000048", "UKAZ": "E12000049", "UKBA": "E12000050", "UKBB": "E12000051", "UKBC": "E12000052", "UKBD": "E12000053", "UKBE": "E12000054", "UKBF": "E12000055", "UKBG": "E12000056", "UKBH": "E12000057", "UKBI": "E12000058", "UKBJ": "E12000059", "UKBK": "E12000060", "UKBL": "E12000061", "UKBM": "E12000062", "UKBN": "E12000063", "UKBO": "E12000064", "UKBP": "E12000065", "UKBQ": "E12000066", "UKBR": "E12000067", "UKBS": "E12000068", "UKBT": "E12000069", "UKBU": "E12000070", "UKBV": "E12000071", "UKBW": "E12000072", "UKBX": "E12000073", "UKBY": "E12000074", "UKBZ": "E12000075", "UKCA": "E12000076", "UKCB": "E12000077", "UKCC": "E12000078", "UKCD": "E12000079", "UKCE": "E12000080", "UKCF": "E12000081", "UKCG": "E12000082", "UKCH": "E12000083", "UKCI": "E12000084", "UKCJ": "E12000085", "UKCK": "E12000086", "UKCL": "E12000087", "UKCM": "E12000088", "UKCN": "E12000089", "UKCO": "E12000090", "UKCP": "E12000091", "UKCQ": "E12000092", "UKCR": "E12000093", "UKCS": "E12000094", "UKCT": "E12000095", "UKCU": "E12000096", "UKCV": "E12000097", "UKCW": "E12000098", "UKCX": "E12000099", "UKCY": "E12000100", "UKCZ": "E12000101", "UKDA": "E12000102", "UKDB": "E12000103", "UKDC": "E12000104", "UKDD": "E12000105", "UKDE": "E12000106", "UKDF": "E12000107", "UKDG": "E12000108", "UKDH": "E12000109", "UKDI": "E12000110", "UKDJ": "E12000111", "UKDK": "E12000112", "UKDL": "E12000113", "UKDM": "E12000114", "UKDN": "E12000115", "UKDO": "E12000116", "UKDP": "E12000117", "UKDQ": "E12000118", "UKDR": "E12000119", "UKDS": "E12000120", "UKDT": "E12000121", "UKDU": "E12000122", "UKDV": "E12000123", "UKDW": "E12000124", "UKDX": "E12000125", "UKDY": "E12000126", "UKDZ": "E12000127", "UKEA": "E12000128", "UKEB": "E12000129", "UKEC": "E12000130", "UKED": "E12000131", "UKEE": "E12000132", "UKEF": "E12000133", "UKEG": "E12000134", "UKEH": "E12000135", "UKEI": "E12000136", "UKEJ": "E12000137", "UKEK": "E12000138", "UKEL": "E12000139", "UKEM": "E12000140", "UKEN": "E12000141", "UKEO": "E12000142", "UKEP": "E12000143", "UKEQ": "E12000144", "UKER": "E12000145", "UKES": "E12000146", "UKET": "E12000147", "UKEU": "E12000148", "UKEV": "E12000149", "UKEW": "E12000150", "UKEX": "E12000151", "UKEY": "E12000152", "UKEZ": "E12000153", "UKFA": "E12000154", "UKFB": "E12000155", "UKFC": "E12000156", "UKFD": "E12000157", "UKFE": "E12000158", "UKFF": "E12000159", "UKFG": "E12000160", "UKFH": "E12000161", "UKFI": "E12000162", "UKFJ": "E12000163", "UKFK": "E12000164", "UKFL": "E12000165", "UKFM": "E12000166", "UKFN": "E12000167", "UKFO": "E12000168", "UKFP": "E12000169", "UKFQ": "E12000170", "UKFR": "E12000171", "UKFS": "E12000172", "UKFT": "E12000173", "UKFU": "E12000174", "UKFV": "E12000175", "UKFW": "E12000176", "UKFX": "E12000177", "UKFY": "E12000178", "UKFZ": "E12000179", "UKGA": "E12000180", "UKGB": "E12000181", "UKGC": "E12000182", "UKGD": "E12000183", "UKGE": "E12000184", "UKGF": "E12000185", "UKGG": "E12000186", "UKGH": "E12000187", "UKGI": "E12000188", "UKGJ": "E12000189", "UKGK": "E12000190", "UKGL": "E12000191", "UKGM": "E12000192", "UKGN": "E12000193", "UKGO": "E12000194", "UKGP": "E12000195", "UKGQ": "E12000196", "UKGR": "E12000197", "UKGS": "E12000198", "UKGT": "E12000199", "UKGU": "E12000200", "UKGV": "E12000201", "UKGW": "E12000202", "UKGX": "E12000203", "UKGY": "E12000204", "UKGZ": "E12000205", "UKHA": "E12000206", "UKHB": "E12000207", "UKHC": "E12000208", "UKHD": "E12000209", "UKHE": "E12000210", "UKHF": "E12000211", "UKHG": "E12000212", "UKHH": "E12000213", "UKHI": "E12000214", "UKHJ": "E12000215", "UKHK": "E12000216", "UKHL": "E12000217", "UKHM": "E12000218", "UKHN": "E12000219", "UKHO": "E12000220", "UKHP": "E12000221", "UKHQ": "E12000222", "UKHR": "E12000223", "UKHS": "E12000224", "UKHT": "E12000225", "UKHU": "E12000226", "UKHV": "E12000227", "UKHW": "E12000228", "UKHX": "E12000229", "UKHY": "E12000230", "UKHZ": "E12000231", "UKIA": "E12000232", "UKIB": "E12000233", "UKIC": "E12000234", "UKID": "E12000235", "UKIE": "E12000236", "UKIF": "E12000237", "UKIG": "E12000238", "UKIH": "E12000239", "UKIJ": "E12000240", "UKIK": "E12000241", "UKIL": "E12000242", "UKIM": "E12000243", "UKIN": "E12000244", "UKIO": "E12000245", "UKIP": "E12000246", "UKIQ": "E12000247", "UKIR": "E12000248", "UKIS": "E12000249", "UKIT": "E12000250", "UKIU": "E12000251", "UKIV": "E12000252", "UKIW": "E12000253", "UKIX": "E12000254", "UKIY": "E12000255", "UKIZ": "E12000256", "UKJA": "E12000257", "UKJB": "E12000258", "UKJC": "E12000259", "UKJD": "E12000260", "UKJE": "E12000261", "UKJF": "E12000262", "UKJG": "E12000263", "UKJH": "E12000264", "UKJI": "E12000265", "UKJJ": "E12000266", "UKJK": "E12000267", "UKJL": "E12000268", "UKJM": "E12000269", "UKJN": "E12000270", "UKJO": "E12000271", "UKJP": "E12000272", "UKJQ": "E12000273", "UKJR": "E12000274", "UKJS": "E12000275", "UKJT": "E12000276", "UKJU": "E12000277", "UKJV": "E12000278", "UKJW": "E12000279", "UKJX": "E12000280", "UKJY": "E12000281", "UKJZ": "E12000282", "UKKA": "E12000283", "UKKB": "E12000284", "UKKC": "E12000285", "UKKD": "E12000286", "UKKE": "E12000287", "UKKF": "E12000288", "UKKG": "E12000289", "UKKH": "E12000290", "UKKI": "E12000291", "UKKJ": "E12000292", "UKKK": "E12000293", "UKKL": "E12000294", "UKKM": "E12000295", "UKKN": "E12000296", "UKKO": "E12000297", "UKKP": "E12000298", "UKKQ": "E12000299", "UKKR": "E12000300", "UKKS": "E12000301", "UKKT": "E12000302", "UKKU": "E12000303", "UKKV": "E12000304", "UKKW": "E12000305", "UKKX": "E12000306", "UKKY": "E12000307", "UKKZ": "E12000308", "UKLA": "E12000309", "UKLB": "E12000310", "UKLC": "E12000311", "UKLD": "E12000312", "UKLE": "E12000313", "UKLF": "E12000314", "UKLG": "E12000315", "UKLH": "E12000316", "UKLI": "E12000317", "UKLJ": "E12000318", "UKLK": "E12000319", "UKLL": "E12000320", "UKLM": "E12000321", "UKLN": "E12000322", "UKLO": "E12000323", "UKLP": "E12000324", "UKLQ": "E12000325", "UKLR": "E12000326", "UKLS": "E12000327", "UKLT": "E12000328", "UKLU": "E12000329", "UKLV": "E12000330", "UKLW": "E12000331", "UKLX": "E12000332", "UKLY": "E12000333", "UKLZ": "E12000334", "UKMA": "E12000335", "UKMB": "E12000336", "UKMC": "E12000337", "UKMD": "E12000338", "UKME": "E12000339", "UKMF": "E12000340", "UKMG": "E12000341", "UKMH": "E12000342", "UKMI": "E12000343", "UKMJ": "E12000344", "UKMK": "E12000345", "UKML": "E12000346", "UKMN": "E12000347", "UKMO": "E12000348", "UKMP": "E12000349", "UKMQ": "E12000350", "UKMR": "E12000351", "UKMS": "E12000352", "UKMT": "E12000353", "UKMU": "E12000354", "UKMV": "E12000355", "UKMW": "E12000356", "UKMX": "E12000357", "UKMY": "E12000358", "UKMZ": "E12000359", "UKNA": "E12000360", "UKNB": "E12000361", "UKNC": "E12000362", "UKND": "E12000363", "UKNE": "E12000364", "UKNF": "E12000365", "UKNG": "E12000366", "UKNH": "E12000367", "UKNI": "E12000368", "UKNJ": "E12000369", "UKNK": "E12000370", "UKNL": "E12000371", "UKNM": "E12000372", "UKNN": "E12000373", "UKNO": "E12000374", "UKNP": "E12000375", "UKNQ": "E12000376", "UKNR": "E12000377", "UKNS": "E12000378", "UKNT": "E12000379", "UKNU": "E12000380", "UKNV": "E12000381", "UKNW": "E12000382", "UKNX": "E12000383", "UKNY": "E12000384", "UKNZ": "E12000385", "UKOA": "E12000386", "UKOB": "E12000387", "UKOC": "E12000388", "UKOD": "E12000389", "UKOE": "E12000390", "UKOF": "E12000391", "UKOG": "E12000392", "UKOH": "E12000393", "UKOI": "E12000394", "UKOJ": "E12000395", "UKOK": "E12000396", "UKOL": "E12000397", "UKOM": "E12000398", "UKON": "E12000399", "UKOO": "E12000400", "UKOP": "E12000401", "UKOQ": "E12000402", "UKOR": "E12000403", "UKOS": "E12000404", "UKOT": "E12000405", "UKOU": "E12000406", "UKOV": "E12000407", "UKOW": "E12000408", "UKOX": "E12000409", "UKOY": "E12000410", "UKOZ": "E12000411", "UKPA": "E12000412", "UKPB": "E12000413", "UKPC": "E12000414", "UKPD": "E12000415", "UKPE": "E12000416", "UKPF": "E12000417", "UKPG": "E12000418", "UKPH": "E12000419", "UKPI": "E12000420", "UKPJ": "E12000421", "UKPK": "E12000422", "UKPL": "E12000423", "UKPM": "E12000424", "UKPN": "E12000425", "UKPO": "E12000426", "UKPP": "E12000427", "UKPQ": "E12000428", "UKPR": "E12000429", "UKPS": "E12000430", "UKPT": "E12000431", "UKPU": "E12000432", "UKPV": "E12000433", "UKPW": "E12000434", "UKPX": "E12000435", "UKPY": "E12000436", "UKPZ": "E12000437", "UKQA": "E12000438", "UKQB": "E12000439", "UKQC": "E12000440", "UKQD": "E12000441", "UKQE": "E12000442", "UKQF": "E12000443", "UKQG": "E12000444", "UKQH": "E12000445", "UKQI": "E12000446", "UKQJ": "E12000447", "UKQK": "E12000448", "UKQL": "E12000449", "UKQM": "E12000450", "UKQN": "E12000451", "UKQO": "E12000452", "UKQP": "E12000453", "UKQQ": "E12000454", "UKQR": "E12000455", "UKQS": "E12000456", "UKQT": "E12000457", "UKQU": "E12000458", "UKQV": "E12000459", "UKQW": "E12000460", "UKQX": "E12000461", "UKQY": "E12000462", "UKQZ": "E12000463", "UKRA": "E12000464", "UKRB": "E12000465", "UKRC": "E12000466", "UKRD": "E12000467", "UKRE": "E12000468", "UKRF": "E12000469", "UKRG": "E12000470", "UKRH": "E12000471", "UKRI": "E12000472", "UKRJ": "E12000473", "UKRK": "E12000474", "UKRL": "E12000475", "UKRM": "E12000476", "UKRN": "E12000477", "UKRO": "E12000478", "UKRP": "E12000479", "UKRQ": "E12000480", "UKRR": "E12000481", "UKRS": "E12000482", "UKRT": "E12000483", "UKRU": "E12000484", "UKRV": "E12000485", "UKRW": "E12000486", "UKRX": "E12000487", "UKRY": "E12000488", "UKRZ": "E12000489", "UKSA": "E12000490", "UKSB": "E12000491", "UKSC": "E12000492", "UKSD": "E12000493", "UKSE": "E12000494", "UKSF": "E12000495", "UKSG": "E12000496", "UKSH": "E12000497", "UKSI": "E12000498", "UKSJ": "E12000499", "UKSK": "E12000500", "UKSL": "E12000501", "UKSM": "E12000502", "UKSN": "E12000503", "UKSO": "E12000504", "UKSP": "E12000505", "UKSQ": "E12000506", "UKSR": "E12000507", "UKSS": "E12000508", "UKST": "E12000509", "UKSU": "E12000510", "UKSV": "E12000511", "UKSW": "E12000512", "UKSX": "E12000513", "UKSY": "E12000514", "UKSZ": "E12000515", "UKTA": "E12000516", "UKTB": "E12000517", "UKTC": "E12000518", "UKTD": "E12000519", "UKTE": "E12000520", "UKTF": "E12000521", "UKTG": "E12000522", "UKTH": "E12000523", "UKTI": "E12000524", "UKTJ": "E12000525", "UKTK": "E12000526", "UKTL": "E12000527", "UKTM": "E12000528", "UKTN": "E12000529", "UKTO": "E12000530", "UKTP": "E12000531", "UKTQ": "E12000532", "UKTR": "E12000533", "UKTS": "E12000534", "UKTT": "E12000535", "UKTU": "E12000536", "UKTV": "E12000537", "UKTW": "E12000538", "UKTX": "E12000539", "UKTY": "E12000540", "UKTZ": "E12000541", "UKUA": "E12000542", "UKUB": "E12000543", "UKUC": "E12000544", "UKUD": "E12000545", "UKUE": "E12000546", "UKUF": "E12000547", "UKUG": "E12000548", "UKUH": "E12000549", "UKUI": "E12000550", "UKUJ": "E12000551", "UKUK": "E12000552", "UKUL": "E12000553", "UKUM": "E12000554", "UKUN": "E12000555", "UKUO": "E12000556", "UKUP": "E12000557", "UKUQ": "E12000558", "UKUR": "E12000559", "UKUS": "E12000560", "UKUT": "E12000561", "UKUU": "E12000562", "UKUV": "E12000563", "UKUW": "E12000564", "UKUX": "E12000565", "UKUY": "E12000566", "UKUZ": "E12000567", "UKVA": "E12000568", "UKVB": "E12000569", "UKVC": "E12000570", "UKVD": "E12000571", "UKVE": "E12000572", "UKVF": "E12000573", "UKVG": "E12000574", "UKVH": "E12000575", "UKVI": "E12000576", "UKVJ": "E12000577", "UKVK": "E12000578", "UKVL": "E12000579", "UKVM": "E12000580", "UKVN": "E12000581", "UKVO": "E12000582", "UKVP": "E12000583", "UKVQ": "E12000584", "UKVR": "E12000585", "UKVS": "E12000586", "UKVT": "E12000587", "UKVU": "E12000588", "UKVV": "E12000589", "UKVW": "E12000590", "UKVX": "E12000591", "UKVY": "E12000592", "UKVZ": "E12000593", "UKWA": "E12000594", "UKWB": "E12000595", "UKWC": "E12000596", "UKWD": "E12000597", "UKWE": "E12000598", "UKWF": "E12000599", "UKWG": "E12000600", "UKWH": "E12000601", "UKWI": "E12000602", "UKWJ": "E12000603", "UKWK": "E12000604", "UKWL": "E12000605", "UKWM": "E12000606", "UKWN": "E12000607", "UKWO": "E12000608", "UKWP": "E12000609", "UKWQ": "E12000610", "UKWR": "E12000611", "UKWS": "E12000612", "UKWT": "E12000613", "UKWU": "E12000614", "UKWV": "E12000615", "UKWW": "E12000616", "UKWX": "E12000617", "UKWY": "E12000618", "UKWZ": "E12000619", "UKXA": "E12000620", "UKXB": "E12000621", "UKXC": "E12000622", "UKXD": "E12000623", "UKXE": "E12000624", "UKXF": "E12000625", "UKXG": "E12000626", "UKXH": "E12000627", "UKXI": "E12000628", "UKXJ": "E12000629", "UKXK": "E12000630", "UKXL": "E12000631", "UKXM": "E12000632", "UKXN": "E12000633", "UKXO": "E12000634", "UKXP": "E12000635", "UKXQ": "E12000636", "UKXR": "E12000637", "UKXS": "E12000638", "UKXT": "E12000639", "UKXU": "E12000640", "UKXV": "E12000641", "UKXW": "E12000642", "UKXX": "E12000643", "UKXY": "E12000644", "UKXZ": "E12000645", "UKYA": "E12000646", "UKYB": "E12000647", "UKYC": "E12000648", "UKYD": "E12000649", "UKYE": "E12000650", "UKYF": "E12000651", "UKYG": "E12000652", "UKYH": "E12000653", "UKYI": "E12000654", "UKYJ": "E12000655", "UKYK": "E12000656", "UKYL": "E12000657", "UKYM": "E12000658", "UKYN": "E12000659", "UKYO": "E12000660", "UKYP": "E12000661", "UKYQ": "E12000662", "UKYR": "E12000663", "UKYS": "E12000664", "UKYT": "E12000665", "UKYU": "E12000666", "UKYV": "E12000667", "UKYW": "E12000668", "UKYX": "E12000669", "UKYY": "E12000670", "UKYZ": "E12000671", "UKZA": "E12000672", "UKZB": "E12000673", "UKZC": "E12000674", "UKZD": "E12000675", "UKZE": "E12000676", "UKZF": "E12000677", "UKZG": "E12000678", "UKZH": "E12000679", "UKZI": "E12000680", "UKZJ": "E12000681", "UKZK": "E12000682", "UKZL": "E12000683", "UKZM": "E12000684", "UKZN": "E12000685", "UKZO": "E12000686", "UKZP": "E12000687", "UKZQ": "E12000688", "UKZR": "E12000689", "UKZS": "E12000690", "UKZT": "E12000691", "UKZU": "E12000692", "UKZV": "E12000693", "UKZW": "E12000694", "UKZX": "E12000695", "UKZY": "E12000696", "UKZZ": "E12000697", "UKAA": "E12000698", "UKAB": "E12000699", "UKAC": "E12000700", "UKAD": "E12000701", "UKAE": "E12000702", "UKAF": "E12000703", "UKAG": "E12000704", "UKAH": "E12000705", "UKAI": "E12000706", "UKAJ": "E12000707", "UKAK": "E12000708", "UKAL": "E12000709", "UKAM": "E12000710", "UKAN": "E12000711", "UKAO": "E12000712", "UKAP": "E12000713", "UKAQ": "E12000714", "UKAR": "E12000715", "UKAS": "E12000716", "UKAT": "E12000717", "UKAU": "E12000718", "UKAV": "E12000719", "UKAW": "E12000720", "UKAX": "E12000721", "UKAY": "E12000722", "UKAZ": "E12000723", "UKBA": "E12000724", "UKBB": "E12000725", "UKBC": "E12000726", "UKBD": "E12000727", "UKBE": "E12000728", "UKBF": "E12000729", "UKBG": "E12000730", "UKBH": "E12000731", "UKBI": "E12000732", "UKBJ": "E12000733", "UKBK": "E12000734", "UKBL": "E12000735", "UKBM": "E12000736", "UKBN": "E12000737", "UKBO": "E12000738", "UKBP": "E12000739", "UKBQ": "E12000740", "UKBR": "E12000741", "UK
```



Out[18]:

	region	record_counts
0	E12000001	26767.0
1	E12000002	75080.0
2	E12000003	55469.0
3	E12000004	49440.0
4	E12000005	60231.0
5	E12000006	64258.0
6	E12000007	89489.0
7	E12000008	94344.0
8	E12000009	57815.0
9	W92000004	31458.0
10	S99999999	NaN
11	N99999999	NaN

Now that we have the data we want combined with the geographic data, we can plot it

```
In [19]: ax = map.plot(
    figsize=(10,6),
    column="record_counts",
    legend=True,
    legend_kwds={"label": "Records Per Region", "orientation": "vertical"},
    missing_kwds={
        "color": "lightgrey",
        "edgecolor": "red",
        "hatch": "///",
        "label": "Missing values",
    },
)

ax.set_axis_off()
```



Through this map, we can see that most of the data comes from the Southeast, which makes sense considering it has the highest population density.

We can use a similar idea to find the most responses for any of the statistics by region

Note: To change which stat you are looking at, change the occupation string and corresponding map to whichever column you'd like to observe

```
In [20]: # Make a table of occupation per region
df['stat'] = df['occupation_10a'].map(labels.occupation_labels)
table = pd.crosstab(df["region"], df["stat"])
# Get rid of does not apply so we get interesting results
table = table.drop(columns=["Does not apply"])
```

table									
stat	Administrative and secretarial occupations	Associate professional and technical occupations	Caring, leisure and other service occupations	Elementary occupations	Managers, directors and senior officials	Process, plant and machine operatives	Professional occupations	Sales and customer service occupations	Skilled trades occupations
region									
E12000001	2001	2034	1901	2639	1522	1637	3011	2088	1984
E12000002	5628	5737	5337	7270	5208	4293	9064	5073	5290
E12000003	3947	4132	3822	5709	3902	3382	6269	3764	4221
E12000004	3690	3721	3448	5384	3719	3363	5739	3291	3654
E12000005	4479	4349	4057	6067	4232	3776	6797	3763	4523
E12000006	5391	5423	4440	5567	5656	3400	8299	3945	4904
E12000007	6134	8649	5138	6854	7942	3215	14080	4883	4674
E12000008	8055	8900	6335	7372	9137	3919	13279	5644	6470
E12000009	4964	4942	4378	5111	5170	2855	7606	3864	4973
W92000004	2352	2311	2503	3053	2107	1899	3759	2194	2636

```
In [21]: #Add the maximum occupation for each region as a new column
max_stats = table.idxmax(axis=1)
map = map.join(max_stats.rename("max_stat"), on="region")
map
```

Out[21]:

	region	nuts118nm	bng_e	bng_n	long	lat	GlobalID	geometry	record_counts	max_stat
0	E12000001	North East (England)	417313	600358	-1.728900	55.297031	628608f5-673c-45c8-9082-f57ed9604a20	MULTIPOLYGON (((397942.931 657535.008, 400000 ...	26767.0	Professional occupations
1	E12000002	North West (England)	350015	506280	-2.772370	54.449451	0413c215-94f7-4f2e-a0fd-37d517177413	MULTIPOLYGON (((357136.203 587050.898, 356937....	75080.0	Professional occupations
2	E12000003	Yorkshire and The Humber	446903	448736	-1.287120	53.932640	d608863c-86bd-4877-b3b8-247794a0825a	POLYGON ((479999.978 517670.16, 481340.496 515...	55469.0	Professional occupations
3	E12000004	East Midlands (England)	477660	322635	-0.849670	52.795719	428c9691-f97b-4fa3-aef8-b3d8c0789cf9	POLYGON ((516022.702 412210.895, 515854.099 41...	49440.0	Professional occupations
4	E12000005	West Midlands (England)	386294	295477	-2.203580	52.556969	a8a5953c-6a23-4fca-b1c7-96d80a29dd95	POLYGON ((409402.997 365710.796, 412633.099 36...	60231.0	Professional occupations
5	E12000006	East of England	571074	263229	0.504146	52.240669	5b0cd82b-5790-499e-92b3-b8f1ab8c601c	MULTIPOLYGON (((599964.098 346659.983, 608692....	64258.0	Professional occupations
6	E12000007	London	517516	178392	-0.308640	51.492271	2ae606b7-b7dc-4d57-a422-84064399cbba	POLYGON ((537543.796 199883.103, 537625.1 1960...	89489.0	Professional occupations
7	E12000008	South East (England)	470062	172924	-0.993110	51.450970	53710b5b-1dfe-4a52-bfb9-a07005f418be	MULTIPOLYGON (((493941.101 254585.705, 493776....	94344.0	Professional occupations
8	E12000009	South West (England)	285015	102567	-3.633430	50.811192	da810a21-072d-49d1-853e-48b75dd73c53	MULTIPOLYGON (((416007.596 246053.797, 417496....	57815.0	Professional occupations
9	W92000004	Wales	263406	242881	-3.994160	52.067410	ec38afc1-b206-49a1-bfaf-48c9836e2d02	MULTIPOLYGON (((314837.003 381391.297, 315637....	31458.0	Professional occupations
10	S99999999	Scotland	277746	700060	-3.970910	56.177429	e25e289f-756d-4e60-8388-b4ece419d1c1	MULTIPOLYGON (((320000.017 976689.515, 320250....	NaN	NaN
11	N99999999	Northern Ireland	86601	535325	-6.854810	54.614941	cdeca05f-0e33-4c34-aaf8-af74e1d91622	MULTIPOLYGON (((116908.519 602925.239, 119760....	NaN	NaN

In [22]:

```
ax = map.plot(
    figsize=(10,6),
    column="max_stat",
    legend=True,
    missing_kwds={
        "color": "lightgrey",
        "edgecolor": "red",
        "hatch": "///",
        "label": "Missing values",
    },
)

ax.set_axis_off()
```

